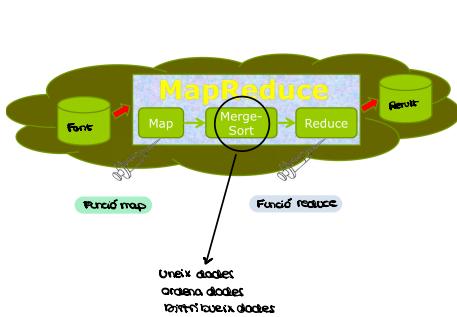


mapReduce

- Treballa sobre Google File System (ara en mes)

- Requirements
 - Tu treballar, ho envies al mapReduce, i dades er distribuixen
 - En comptes de per data shippin, envies query shipping \rightsquigarrow s'executa en local en són les dades
 - Parallelisme:
 - Problema anterior \rightarrow claus parafanes (per fer validacions, joins, ...) \rightarrow dades estan lluny.
 - Vullen treballar amb petabytes.
 - Escalable
 - Si una màquina falla \rightsquigarrow sistema es recupera

- Basat en functional programming



- Processen paralles [key, value]

- Signatura \rightsquigarrow el que hem de posar resultats

`map (key k, value v) \rightarrow [(i1ki, i1vi), ..., (imki, imvi)]`

`reduce (key ik, vset) \rightarrow [ov1, ..., ovr (ik, vs)]`

set
de
values

Exemple: word count

Project Gutenberg Book of the Outline of Science, Vol. 1 (of 4), by J. Arthur Thomson
This eBook is for the use of anyone anywhere at no cost and with almost无限的 freedom. You may copy it, give it away or
re-use it under the terms of the Project Gutenberg License included
with this eBook or online at www.gutenberg.org

Title: The Outline of Science, Vol. 1 (of 4)
A Plain Story Simply Told
Author: J. Arthur Thomson
Release Date: January 22, 2007 [EBook #20417]
Language: English
Character set encoding: ASCII
*** START OF THIS PROJECT GUTENBERG EBOOK OUTLINE OF SCIENCE ***

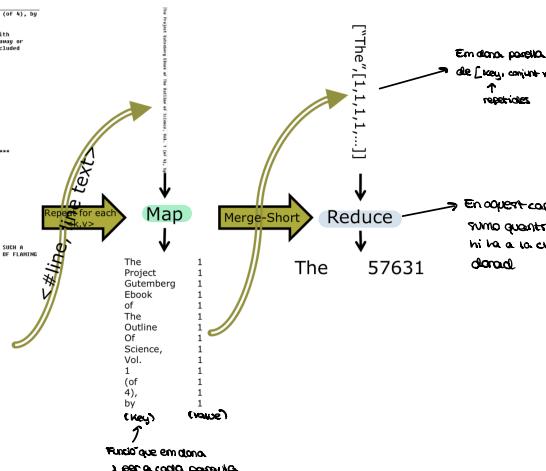
Produced by Brian Jones, Leonard Johnson and the Online Distributed Proofreading Team at http://www.pgdp.net

[Illustration: THE GREAT CORAL REEF PROBLEMS, WHICH ARE SUCH A
MIRACULOUS FEATURE OF THE SOLAR SYSTEM, AND WHICH INDICATE THAT THE
HYDROGEN RISING SOMETHES TO A HEIGHT OF 500,000 MILES]

THE
OUTLINE OF SCIENCE
A PLAIN STORY SIMPLY TOLD

BASIC BY
J. ARTHUR THOMSON
RECEIVED PROFESSOR OF NATURAL HISTORY IN THE
UNIVERSITY OF AMERICA

WITH OVER 400 ILLUSTRATIONS
OF WHICH ABOUT 40 ARE IN COLOR



```
public void map(LongWritable key, Text value) {
    String line = value.toString();
    StringTokenizer tokenizer = new StringTokenizer(line);
    while (tokenizer.hasMoreTokens()) {
        write(new Text(tokenizer.nextToken()), new IntWritable(1));
    }
}

public void reduce(Text key, Iterable<IntWritable> values) {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get();
    }
    write(key, new IntWritable(sum));
}
```

PROBLEMA.

Com que no sap què hi ha al "whatever", no pot optimitzar res.

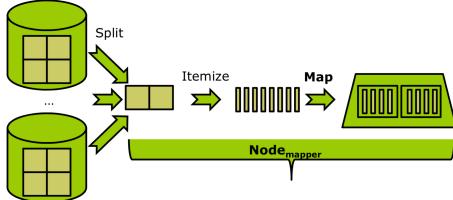
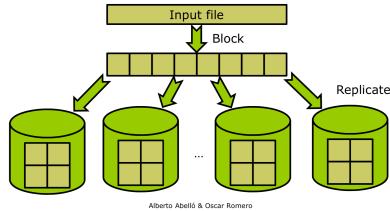
Versió 2 \rightsquigarrow Spark: intenta mirar el "whatever" per optimitzar uns mica.

```
public void map (key, value) {
    whatever
    key, value
}
```

```
public void reduce (key, value) {
    whatever
    key, value
}
```

Algoritme

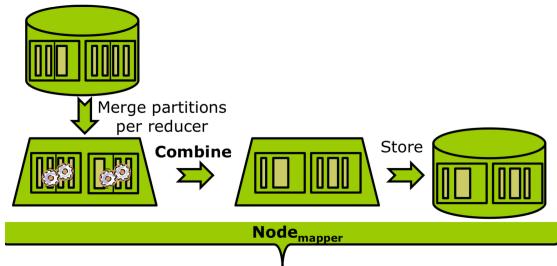
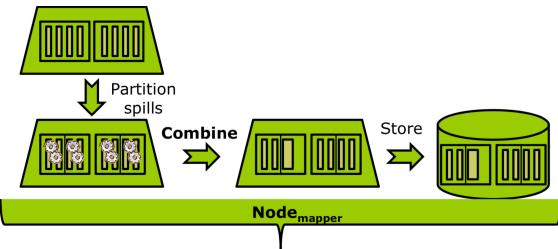
- ② Pujar fitxer a HDFS (cloud)
 - * HDFS ho divideix en chunks
 - * Per defecte, ho replica a 3 instances.



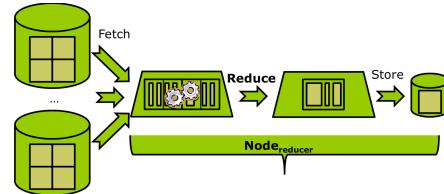
- ③ A cada un dels workers, se li assigna un split (subset de blocks)
- ④ Worker el divideix en records
- ⑤ Executa la funció del map i els posa a memòria (spill)

↳ conte resultats d'haver executat la meia copia "map"

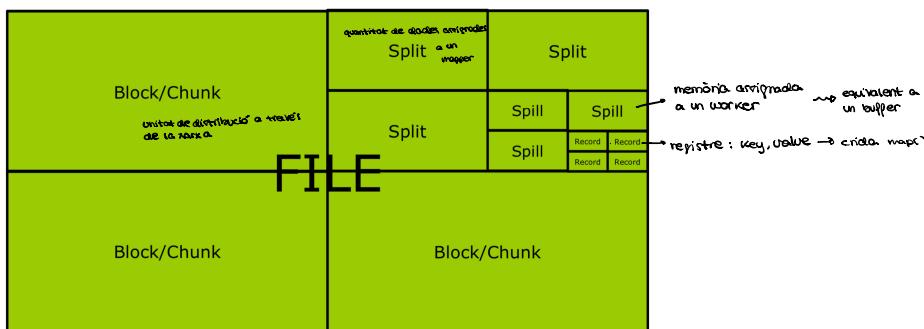
- ⑥ En cada spill, es particiona amb tants trastos com reducers tenim
- ⑦ Cada partió de "spill" s'ordena de forma independent.
- ⑧ Cada cop que ordenem, s'executa "combine"
 - ↑
 - reduce parcial → pel merge.
- ⑨ Guardem les particions de spill a disc (massive writing)



- ⑩ Agafa les particions de spill ordenades i les combina i ordena
- ⑪ Guardem el resultat al disc.



Objectes



Record=Key-Value pair

Combiner

- Permet optimitzar (\downarrow lectures de disc)
- Redueix transit de xarxa
- Important \rightarrow commutativa \lrcorner \rightarrow igual que el reducer
 \hookrightarrow associativa
- $\frac{\text{input}}{\text{output}} \gg \# \text{CPU}$

Problemes mapReduce

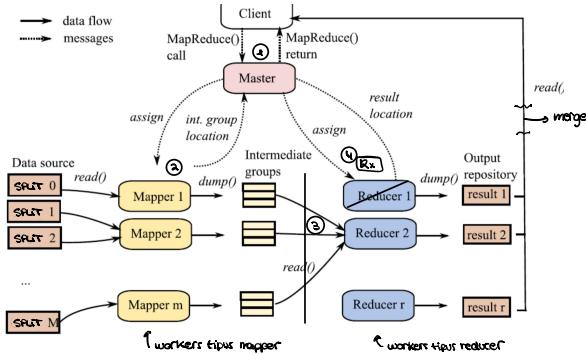
- Massa baix nivell (java) \rightarrow tant eficient com tu el fagis

Avantatges mapReduce

- Pots fer el que vulquis
- Tanca: pots fer cicles entre map-reduce



Tasques i dataflow



PROBLEMES \rightarrow Reassignments needed (4)
 \hookrightarrow Single-point failure (1 master)

Sempre que s'escriu, es fa en local
 Es el "reader" qui va a buscar les dades

| no perdrem
fitxer mai

Costos

- (1) Start-up time
- (2) Writing intermediate results
 \hookrightarrow necessari per no perdre dades)
- (3) Transferència de dades (xarxa)
- (4) Fer pings a master per dir que estan vius
 \hookrightarrow necessari per tenir availability)

Què fa Spark?

- menys escriptures a disc: només quan ho volem necessari
- permet més operacions
- et facilita controlar les particions

Exercici mapreduce

Aixomim que tenim el següent codi:

```
public void map(LongWritable key, Text value) {
    String line = value.toString();
    StringTokenizer tokenizer = new StringTokenizer(line);
    while (tokenizer.hasMoreTokens()) {
        write(new Text(tokenizer.nextToken()), new IntWritable(1));
    }
}

public void reduce(Text key, Iterable<IntWritable> values) {
    int sum = 0;
    for (IntWritable val : values) {
        sum += val.get();
    }
    write(key, new IntWritable(sum));
}
```

#split ← #màquines virtuals

to set:
 1a. Block0: "a b b^c | c d c a e"
 1b. Block1: "a b a d a | b b c c p"

configuration

- combine = reduce
- 1 split = 1 block
- '1' divideix els records del block mitjans de records/block
- dfn. replication = 1 (hadoop)
- 1 spill = 4 [key,value]
- 2 mappers, 2 reducers

↳ machine0 → blocks , machine1 → blocks
 ↳ mapper0 ↳ reducer0 ↳ mapper1 ↳ reducer1

- hash for reducers → {b,a,p} → 0
 ↳ {a,c,e} → 1

Fill the gaps:

- Machine0 contains 1a. blocks. (block0)
 Machine1 contains 1b. blocks. (block1)
- We keep ... replicas (including the master copy) per block.
- We have ... splits per machine. (1split = 1block)
- Mapper0 reads 2.... records.
 Mapper1 reads 2.... records.

- Spills in Machine0:

Spill1	Spill2	Spill3	Spill4
[a,1] [b,1]	[c,1] [d,1]	[e,1] [f,1]	[,] [,]
[b,1] [a,1]	[d,1] [c,1]	[,] [,]	[,] [,]

Spills in Machine1:

Spill1	Spill2	Spill3	Spill4
[a,1] [b,1]	[a,1] [b,1]	[c,1] [f,1]	[,] [,]
[d,1] [a,1]	[b,1] [c,1]	[,] [,]	[,] [,]

arriba vora.

- Partitions in machine0:

Spill 1		Spill 2		Spill 3	
Partition 0	Partition 1	Partition 0	Partition 1	Partition 0	Partition 1
[b,1] [b,1]	[a,1] [a,1]	[d,1] [,]	[e,1] [c,1]	[,] [,]	[a,1] [e,1]
[,] [,]	[,] [,]	[,] [,]	[c,1] [,]	[,] [,]	[,] [,]

Partitions in Machine1:

Spill 1		Spill 2		Spill 3	
Partition 0	Partition 1	Partition 0	Partition 1	Partition 0	Partition 1
[b,1] [d,1]	[a,1] [,]	[b,1] [b,1]	[a,1] [c,1]	[f,1] [,]	[c,1] [,]
[a,1] [,]	[,] [,]	[,] [,]	[,] [,]	[,] [,]	[,] [,]

7) Partitions in machine0:

Spill 1		Spill 2		Spill 3	
Partition 0	Partition 1	Partition 0	Partition 1	Partition 0	Partition 1
[b,2][,]	[a,2][,]	[a,1][,]	[c,2][,]	[,][,]	[a,1][e,1][,]
[,][,]	[,][,]	[,][,]	[,][,]	[,][,]	[,][,]

cada partició
en un spill

← TEORIA
COMBINAR
CONE SUMA

+ ordenar

Partitions in Machine1:

Spill 1		Spill 2		Spill 3	
Partition 0	Partition 1	Partition 0	Partition 1	Partition 0	Partition 1
[b,1][d,2]	[a,2][,]	[b,2][,]	[a,1][c,1]	[g,1][,]	[c,1][,]
[,][,]	[,][,]	[,][,]	[,][,]	[,][,]	[,][,]

8) Files in machine0:

File0	File1	File2
[b,2][,][,][,]	[a,2][,][,][,]	[d,1][,][,][,]
File3	File4	File5
[c,3][,][,][,]	[,][,][,][,]	[a,1][e,1][,][,]
Files in Machine1:		
File0	File1	File2
[b,1][d,2][,][,]	[a,1][,][,][,]	[b,2][,][,][,]
File3	File4	File5
[a,1][c,1][,][,]	[g,1][,][,][,]	[c,1][,][,][,]

9) Merges in machine0:

Merge0	Merge1
[b,2][d,1][,][,][,]	[a,2][c,3][e,1][,][,][,]
Merges in Machine1:	
Merge0	Merge1
[b,3][d,2][g,1][,][,][,]	[a,2][c,2][,][,][,][,]

10) Files in machine0:

File0	File1
[b,2][d,1][,][,][,]	[a,3][c,3][e,1][,][,][,]
Files in Machine1:	
File0	File1
[b,3][d,2][g,1][,][,][,]	[a,2][c,2][,][,][,][,]

11) Reducer0 reads ... File 0 ... Files from machine0

and ... File 0 ... Files from machine1. (answer which Files)

Reducer1 reads ... File 1 ... Files from machine0
and ... File 1 ... Files from machine1. (answer which Files)

12) Merge in machine0:

File0	File1
[b,12,34][d,4,1,2,8][g,1,2,4]	[,][,][,][,][,]
Merge in Machine1:	
File0	File1
[a,43,14][e,4,3,2,4][c,1,1]	[,][,][,][,][,]

13) Reduce function is executed ... times in machine0.

Reduce function is executed ... times in machine1.

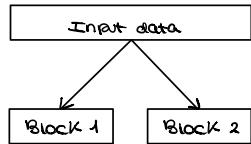
14) Files in machine0:

File0	File1
[b,5][d,3][g,1][,][,][,]	[,][,][,][,][,]
Files in Machine1:	
File0	File1
[a,4][c,5][e,4][,][,][,]	[,][,][,][,][,]

Algorithme pas à pas.

- ① Input data partitioned into blocks

Répartir les données au cloud par blocks \rightarrow Block0
Block1



- ② Replicate them into different nodes

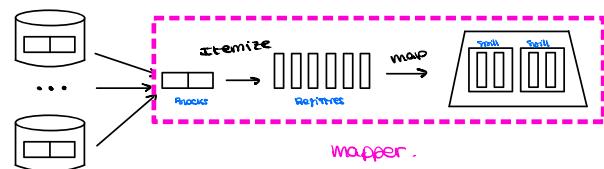
2 replicas/mot



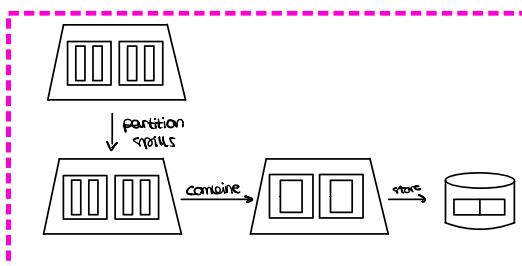
- ③ Each map function reads a subset of blocks (split)

1 mapper \rightarrow 1 block = 1 split

- ④ Divides into records. ~ 2 records / block (split)



- ⑤ Executes the map for each record. ~ 4 records \rightarrow 4 maps \rightarrow on pair segs map



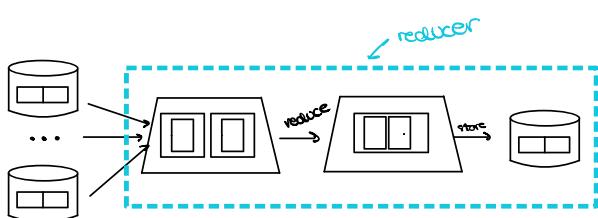
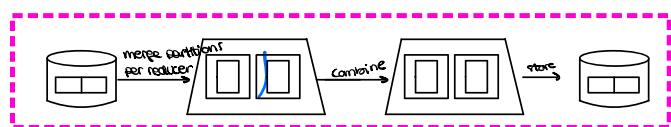
- ⑥ Each spill is partitioned per reducers \rightarrow 2 reducers \rightarrow 2 spills

- ⑦ Each spill is stored independently (+combine)

- ⑧ Store spill partitions into disk

- ⑨ Spill partitions are merged and stored independently (+combine)

- ⑩ store into disk.



- ⑪ Reducers fetch data from mappers

- ⑫ Mapper output is merged and stored

- ⑬ Reduce function is executed per key

- ⑭ Store the result into disk.