

---

# Principles of Distributed Databases

# Knowledge Objectives

---

1. Name the two main characteristics of a distributed database
2. Explain what are the benefits of data fragmentation and data replication performance-wise
3. Explain the 7 degrees of transparency a DDBMS might provide
4. Enumerate the 4 kinds of distributed databases attending to their autonomy
5. Justify the need of new schema levels in the ANSI-SPARC architecture due to data distribution
6. Enumerate and justify the need of new components in the functional architecture of a DDBMS
7. Enumerate the contents of the global catalog
8. Explain 7 bottlenecks of traditional relational systems
9. Explain 6 simplifications for RDBMS in an OLTP environment
10. Enumerate 6 environments that entail different access (and consequently storage) characteristics
11. Explain what are the main goals behind the NOSQL wave
12. Enumerate 4 challenges in data distribution database

# Understanding Objectives

---


1. Compare the latency and response time of a query in both centralized and distributed storage

## Principles of Distributed Databases

# MOTIVATION

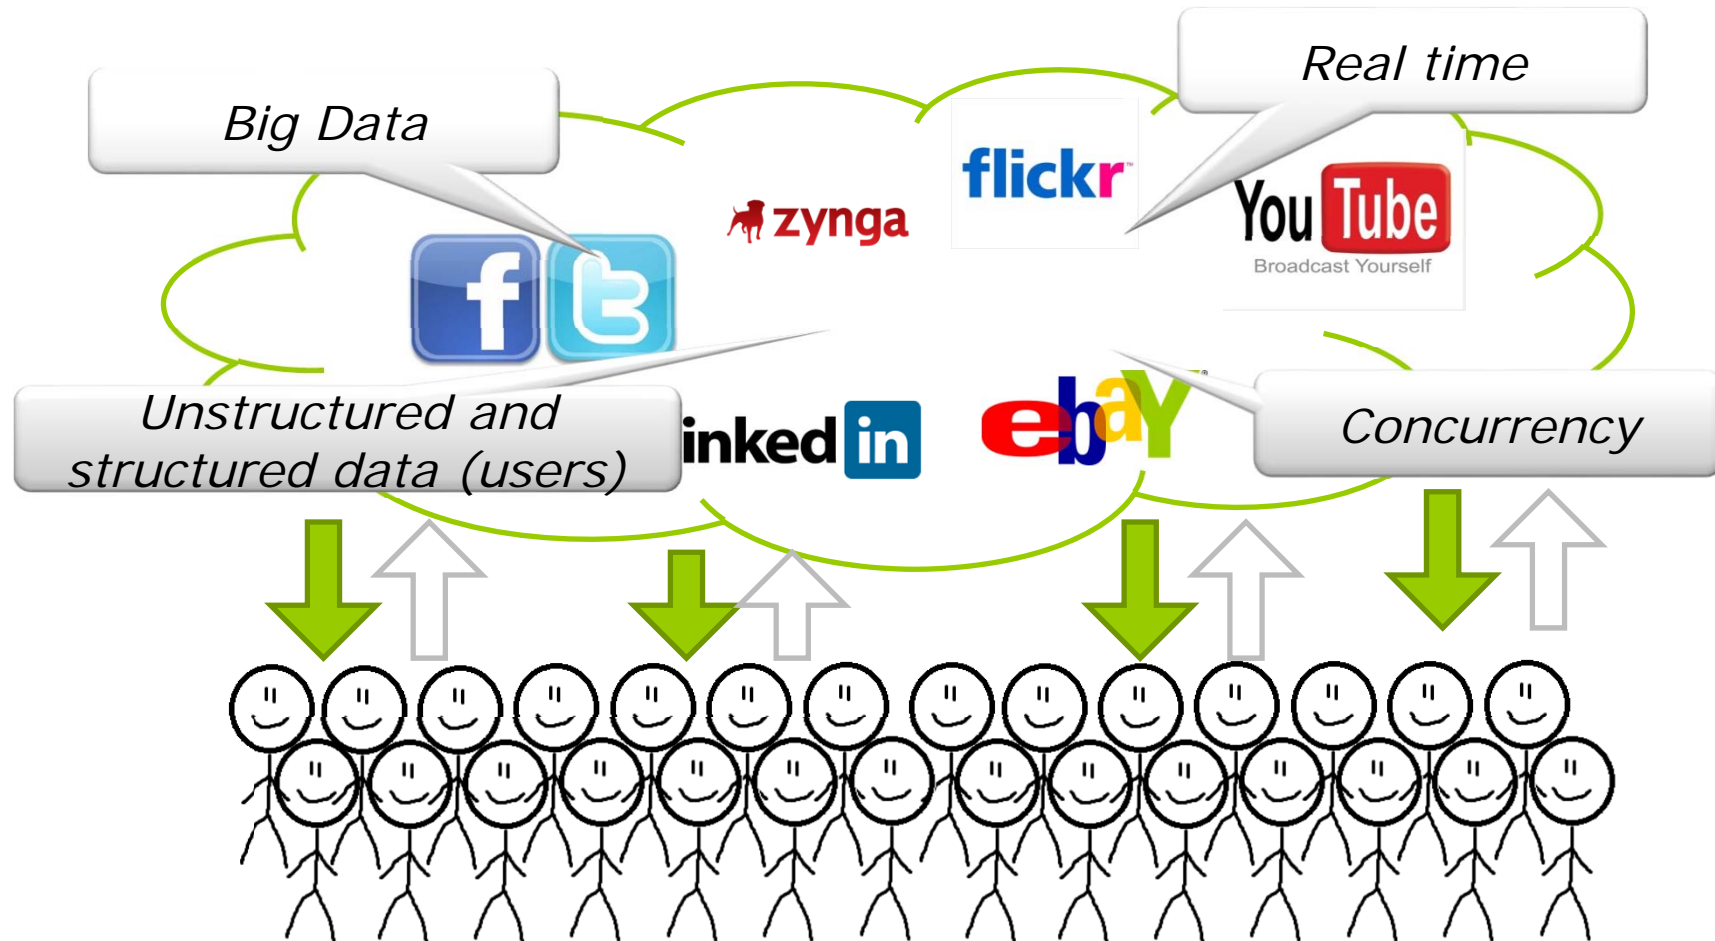
# The End of an Architectural Era

## WEB 1.0 – Read Era



# The End of an Architectural Era

## WEB 2.0 – Write Era



# RDBMS: One Size Does Fit All


- Mainly write-only Systems (e.g., OLTP)
  - Data storage
    - Normalization
  - Queries
    - Indexes: B+, Hash
    - Joins: BNL, RNL, Hash Join, Merge Join
- Read-only Systems (e.g., DW)
  - Data Storage
    - Denormalized data
  - Queries
    - Indexes: Bitmaps
    - Joins: Star-join
    - Materialized Views

**BUT, WHAT IF I NEED TO  
DEAL WITH MASSIVE  
READS AND WRITES AT  
THE SAME TIME?**

# RDBMS Approach

---

Too many reads?  data replication

Too many writes?  data fragmentation

- ❑ Distributed RDBMS (DRDBMS) are not flexible enough
  - ACID properties must be preserved when synchronizing nodes
  - Too much logging writes
  - They do not scale well



Principles of Distributed Databases

# DISTRIBUTED DATABASES

# Activity: Why Distribution?


- **Objective:** Recognize the benefits of distributing data
- **Tasks:**

1. (15') By pairs, answer the following questions:
  - a) How long would it take to read 1TB with sequential access (fig. a)? (in secs)
    - Can you identify any additional drawback to be considered?
  - b) How long would it take to read 1TB with parallel access (fig. b)? Assume 100 disks on the same machine with shared-memory and infinite CPU capacity.
    - Can you identify any additional drawback to be considered?
  - c) How long would it take to read 1TB with distributed access (fig. c)? Assume 100 shared-nothing machines in a star-shape LAN in a single rack where all data is sent to the center.
    - Can you identify any additional drawback to be considered?
  - d) Now, repeat the exercise considering a single random access. What changes?
2. (5') Discussion


Type	Latency	Bandwidth
Disk	$\approx 5 \times 10^{-3}$ s (5 millisecc.);	At best 100 MB/s
LAN	$\approx 1 - 2 \times 10^{-3}$ s (1-2 millisecc.);	$\approx 1$ GB/s (single rack); $\approx 100$ MB/s (switched);
Internet	Highly variable. Typ. 10-100 ms.;	Highly variable. Typ. a few MB/s.;

Bottom line (1): it is approx. one order of magnitude faster to exchange main memory data between 2 machines in a data center, that to read on the disk.


Bottom line (2): exchanging through the Internet is slow and unreliable with respect to LANs.



a. Single CPU, single disk



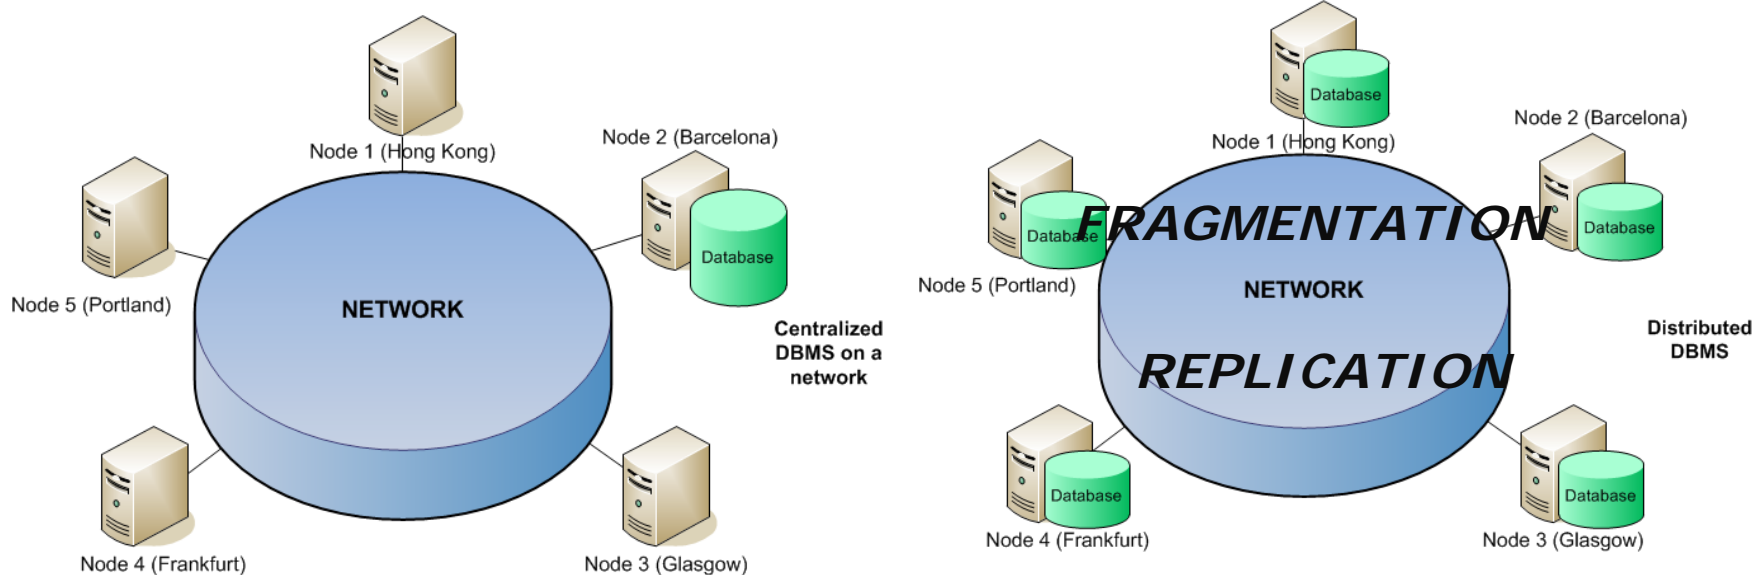
b. Parallel read: single CPU, many disks



c. Distributed reads: an extendible set of servers

# Distributed Database

- ❑ A distributed database (DDB) is a database where data management is distributed over several nodes in a network.
  - Each node is a database itself
    - ❑ Potential heterogeneity
  - Nodes communicate through the network



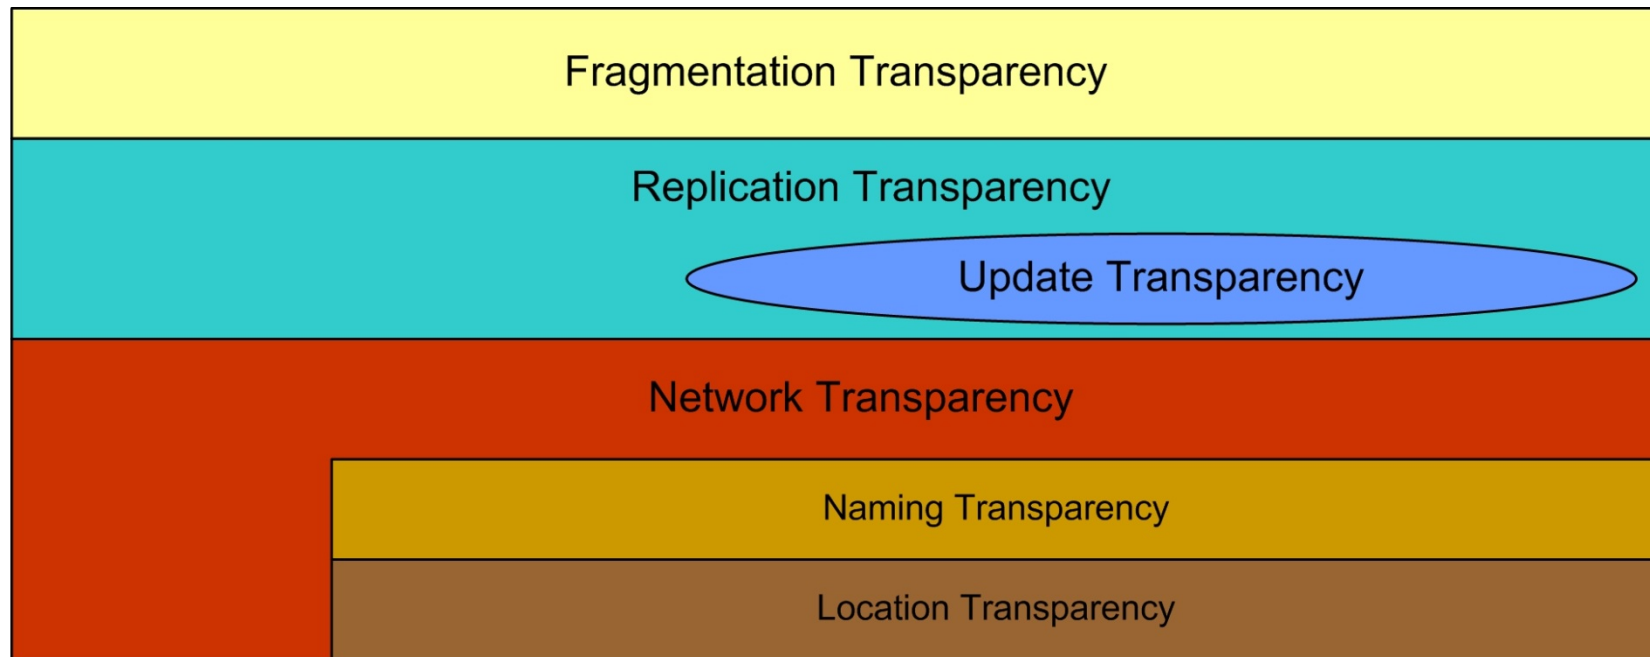
# Distributed Architectures

---

- ❑ Main objective: hide implementation (i.e., physical) details to the users
  - Data independency at the logical and physical level must be guaranteed
    - ❑ Inherited from centralized DBs (ANSI SPARC)
  - Network transparency
    - ❑ Data access must be independent regardless where data is stored
    - ❑ Each data object must have a unique name
  - Replication transparency
    - ❑ The user must not be aware of the existing replicas
  - Fragmentation transparency
    - ❑ The user must not be aware of partitioning

# Distributed Architectures


---



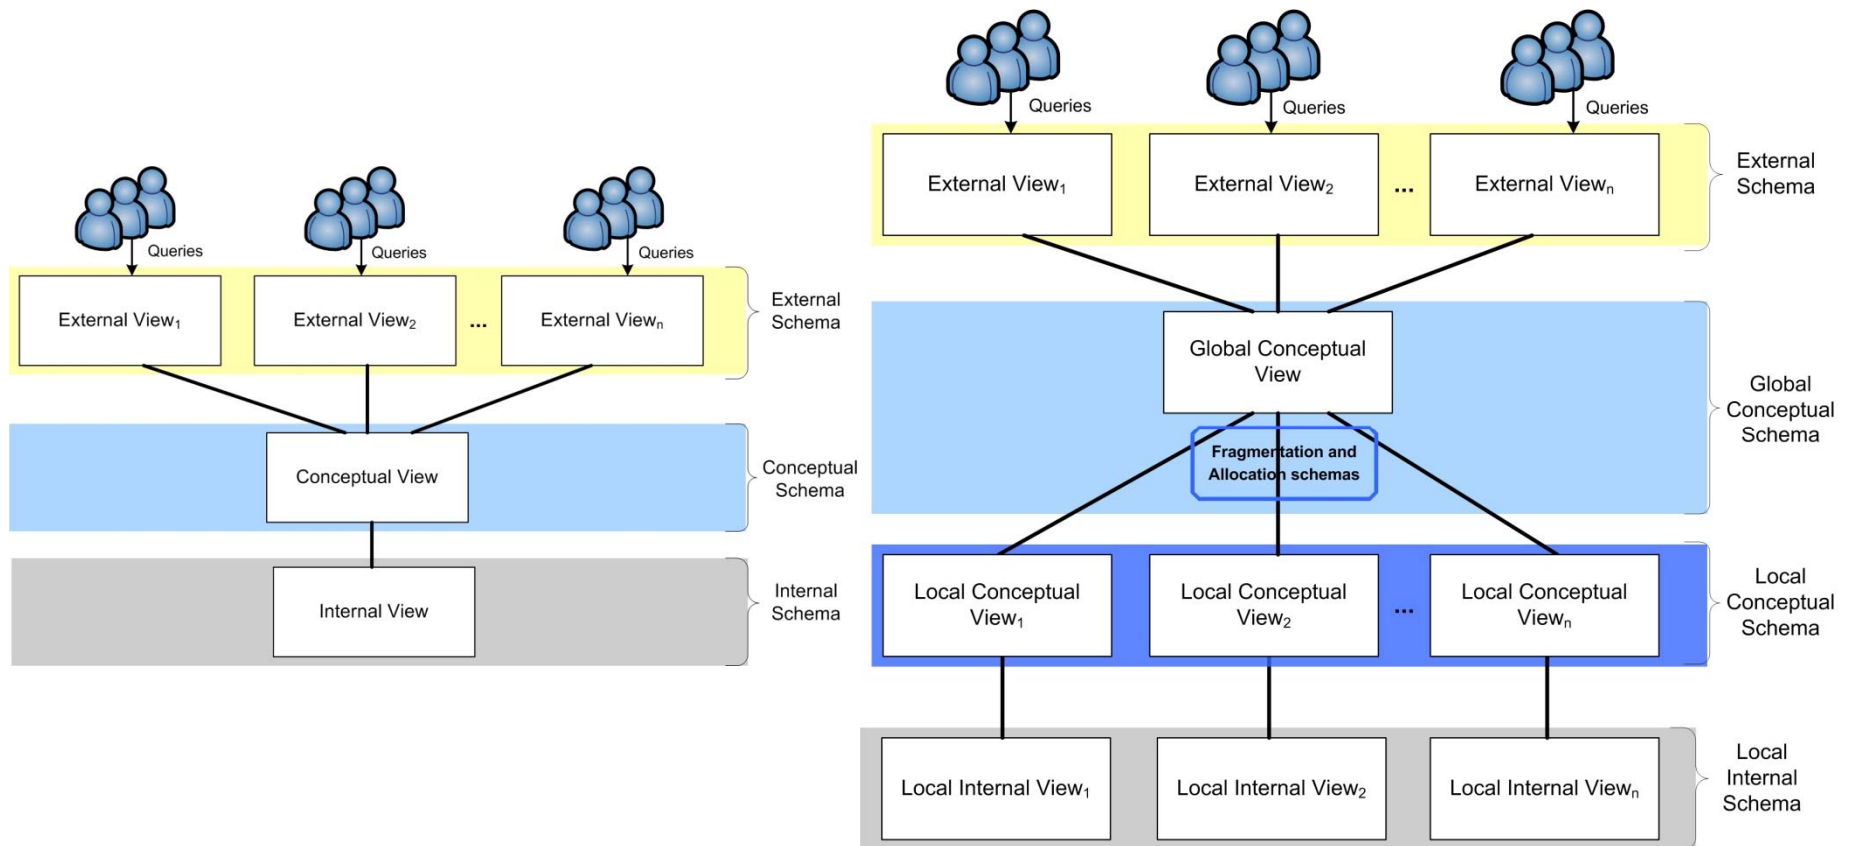
# Comparison of DDBs

---

	<b>Autonomy</b>	<b>Central Schema</b>	<b>Query Transparency</b>	<b>Update Transparency</b>
<i>Homogeneous DBs</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>
<i>Tightly Coupled Federated DBs</i>	<i>Low</i>	<i>Yes</i>	<i>Yes</i>	<i>Limited</i>
<i>Loosely Coupled Federated DBs</i>	<i>Medium</i>	<i>No</i>	<i>Yes</i>	<i>Limited</i>
<i>Multi-databases</i>	<i>High</i>	<i>No</i>	<i>No</i>	<i>No</i>

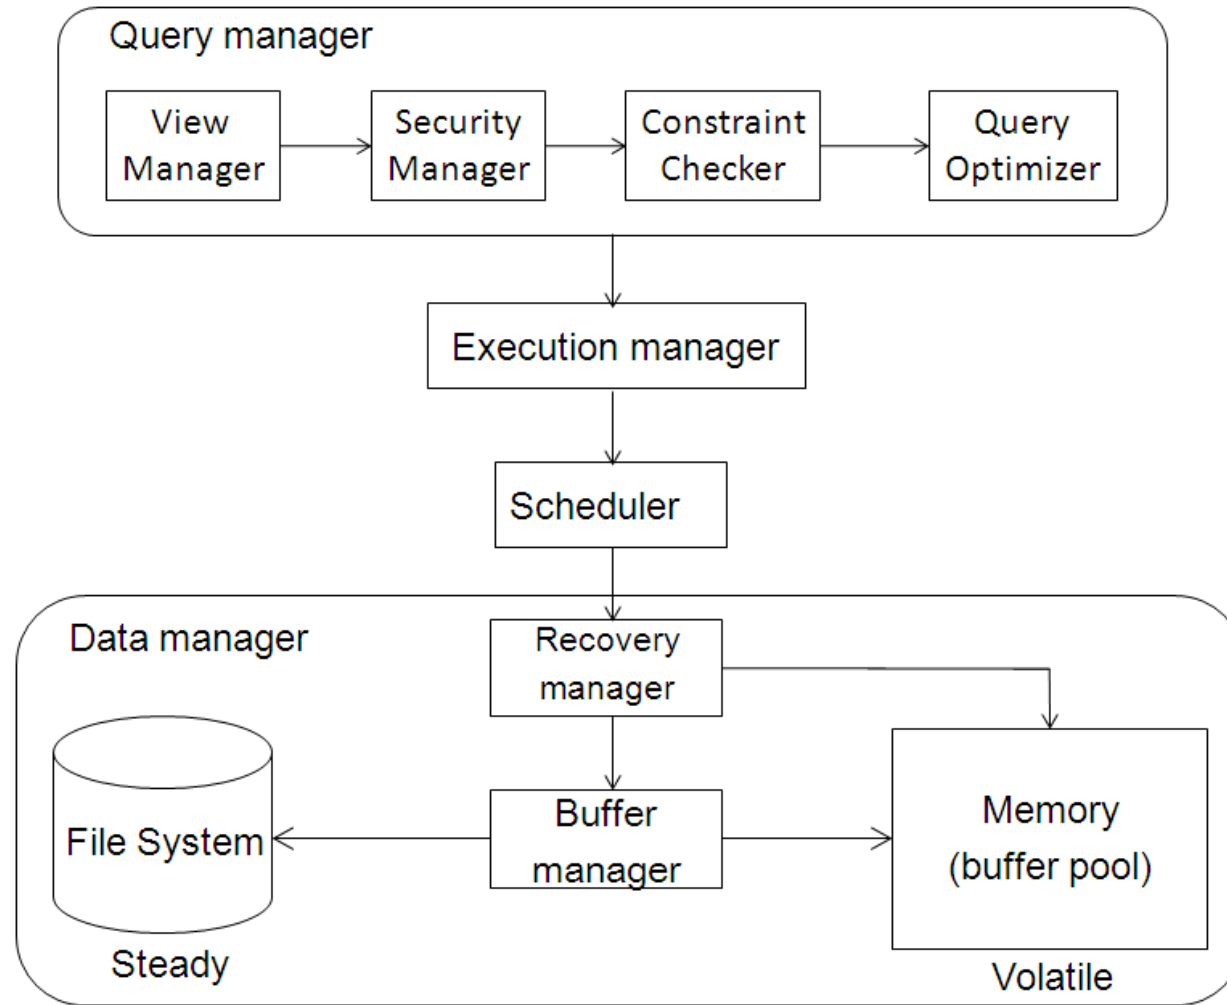


# Extended ANSI-SPARC Architecture



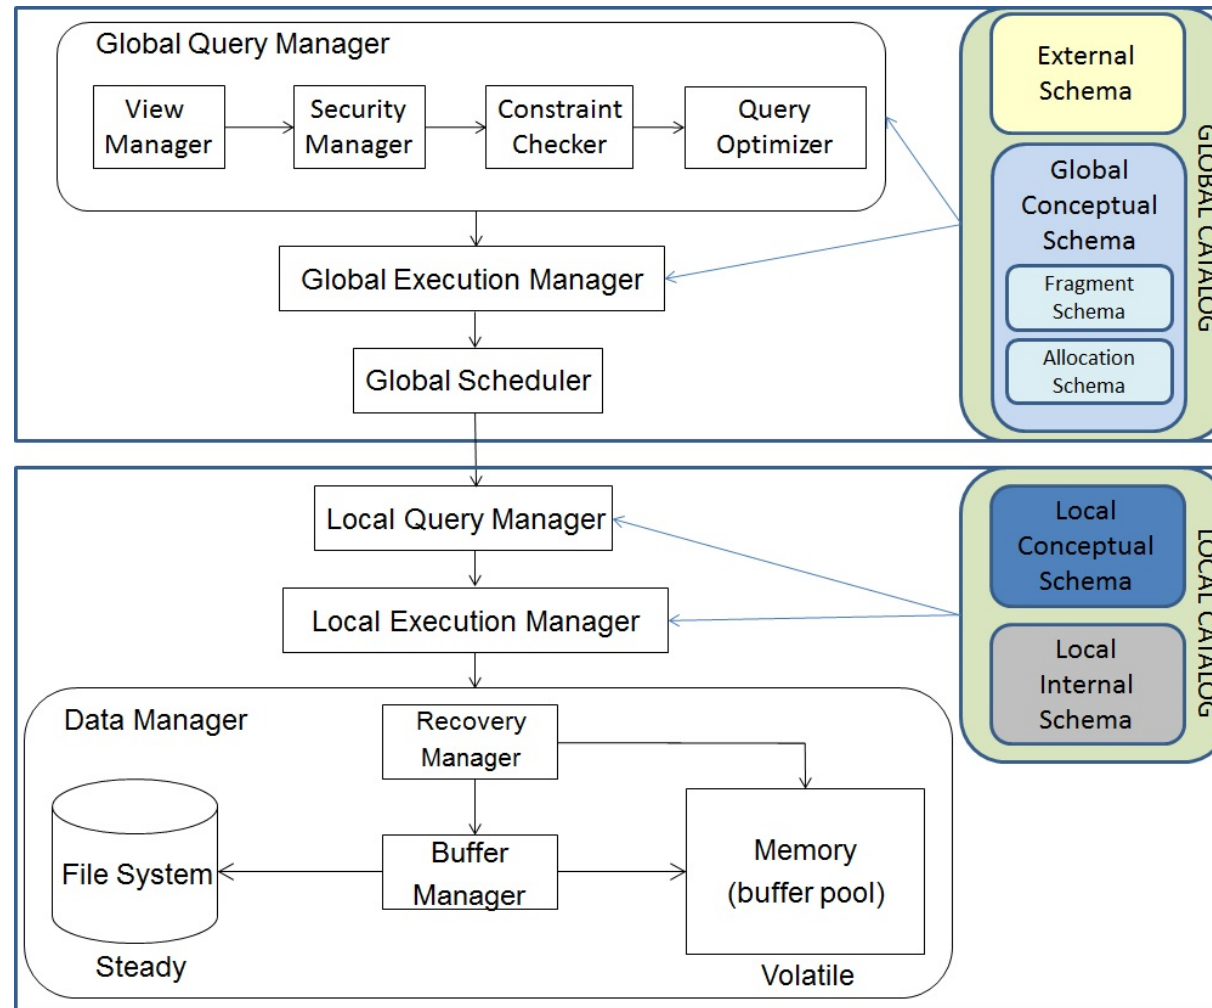
- Global catalog
  - Mappings between ESs – GCS and GCS – LCSs
- Each node has a local catalog
  - Mappings between  $LCS_i$  –  $IS_i$

# Centralized DBMS Architecture





# Distributed DBMS Architecture



## Principles of Distributed Databases

# NOSQL

# Different applications

---

## Not Only SQL (different problems entail different solutions)

- OLTP
  - Object-Relational
    - Distributed databases
    - Parallel databases
- Scientific databases and other Big Data repositories
  - Key-value stores
- Data Warehousing & OLAP
  - MOLAP
  - Column stores
  - Multidimensional features
- Text / documents
  - Document databases
    - XML/JSON databases
- Stream processing
  - Stream processor
- Semantic Web and Open Data
  - Graph databases

# The Problem is Not SQL

---

- Relational systems are too generic...
  - OLTP: stored procedures and simple queries
  - OLAP: ad-hoc complex queries
  - Documents: large objects
  - Streams: time windows with volatile data
  - Scientific: uncertainty and heterogeneity
- ... But the overhead of RDBMS has nothing to do with SQL
  - Low-level, record-at-a-time interface is not the solution

*SQL Databases vS. NoSQL Databases*

*Michael Stonebraker*

*Communications of the ACM, 53(4), 2010*

# Activity: RDBMS Bottlenecks

---

- *Objective: Identify RDBMSs main flaws when used to implement massive distributed systems*
  - *Thus, assume main memory storage, built-in high availability, no user stalls, and useful transaction work under 1 millisecond*
- *Tasks:*
  1. *(5') Read the bottlenecks assigned to you:*
    - i. *Logging*
    - ii. *JDBC connection and latching*
    - iii. *Concurrency control and distributed commit protocols*
  2. *(15') In group of 3 discuss:*
    1. *Each of you start explaining the bottlenecks assigned to you assuming the kind of massive distributed system described above*
    2. *Now, try to generalize them (do not assume a massive distributed system anymore). Can I always get rid of these bottlenecks? In which scenarios could I? In which I could not? Discuss it in terms of storage, running time for transactions and system availability*
  3. *(10') Think tank*
- *Roles for the team-mates during task 2:*
  - a) *Explains his/her material (one at a time)*
  - b) *Asks for clarifications, raises doubts, answers doubts (all)*
  - c) *Mediates and controls time (one at a time)*

# RDBMS Bottlenecks

- ❑ Buffers management (cache disk pages)
- ❑ Logging (WALP)
  - Persistent redo log
  - Undo log
- ❑ Concurrency control (locking)
- ❑ Latching for multi-threading
- ❑ CLI interfaces (JDBC, ODBC, etc.)
- ❑ Variable length records management
  - Locate records in a page
  - Locate fields in a record
- ❑ Two-phase commit protocol in distributed transactions



# NewSQL: A New Architecture for OLTP

**But even for distributed, massive OLTP systems RDBMS can be outperformed!**

- ❑ Main memory DB
  - A DB less than 1Tb fits in memory
    - ❑ 20 nodes x 32 Gb (or more) costs less than 50,000US\$
  - Undo log is in-memory and discarded on commit
- ❑ One thread systems
  - Perform incoming SQL commands to completion, without interruption
    - ❑ One transaction takes less than 1ms
  - No isolation needed
- ❑ Grid computing
  - Enjoy horizontal partitioning and parallelism
    - ❑ Add new nodes to the grid without going down
- ❑ High availability
  - Cannot wait for the recovery process
    - ❑ Multiple machines in a PeerToPeer configuration
- ❑ Reduce costs
  - Human costs are higher than Hw and Sw
    - ❑ An expert DBA is expensive and rare
  - Alternative is brute force
    - ❑ Automatic horizontal partitioning and replication
      - Execute queries at any replica and updates to all of them
- ❑ Optimize queries at compile time

# Summary

---

- ❑ What is a distributed DBMS
- ❑ Distributed architecture for a DDBMS
  - Distribution transparency
  - Replication transparency
  - Fragmentation transparency
- ❑ Benefits of distributed systems
- ❑ NOSQL main goals and features



# Bibliography

---

- ▣ M.T. Özsu and P. Valduriez. *Principles of Distributed Database Systems*. Second edition. Prentice Hall, 1999
- ▣ Serge Abiteboul, Ioana Manolescu, Philippe Rigaux, Marie-Christine Rousset, Pierre Senellart. *Web Data Management*. Cambridge Press, 2011.
- ▣ L. Liu, M.T. Özsu (Eds.). *Encyclopedia of Database Systems*. Springer, 2009

# Recommended Read

---

## ❑ **SQL Databases vS. NoSQL Databases**

Michael Stonebraker

Communications of the ACM, 53(4)

April 2010

- ❑ This paper discusses the RDBMS bottlenecks and underlines the fact that NOSQL is indeed an inappropriate term.
  - This is the main reason why, in this course, you will see NOSQL everywhere instead of NoSQL. NOSQL is not the SQL (or RDBMSs) negation but a complementary view meeting the *one size does not fit all* principle